

[1] "Getting started in Bayesian
modelling with Stan and rstan"

> bill.dixon at alumni.unimelb.edu.au

$$Y = a + b^*X + \varepsilon$$

> ?uncertainty

Aleatory:

- true 'random' variability,
- irreducible but often better characterised

Incertitude / Epistemic:

- Model uncertainty
- abstraction, ignorance
- Often ignored

$$Y \sim \text{normal}(\mu, \sigma)$$

$$\mu = a + b \cdot X$$

> ?bayesian

Diagram illustrating Bayes' theorem:

$$p(\theta|y) = \frac{p(\theta)p(y|\theta)}{p(y)}$$

Labels and arrows:

- posterior points to $p(\theta|y)$
- prior points to $p(\theta)$
- likelihood points to $p(y|\theta)$
- average likelihood / normalising constant points to $p(y)$

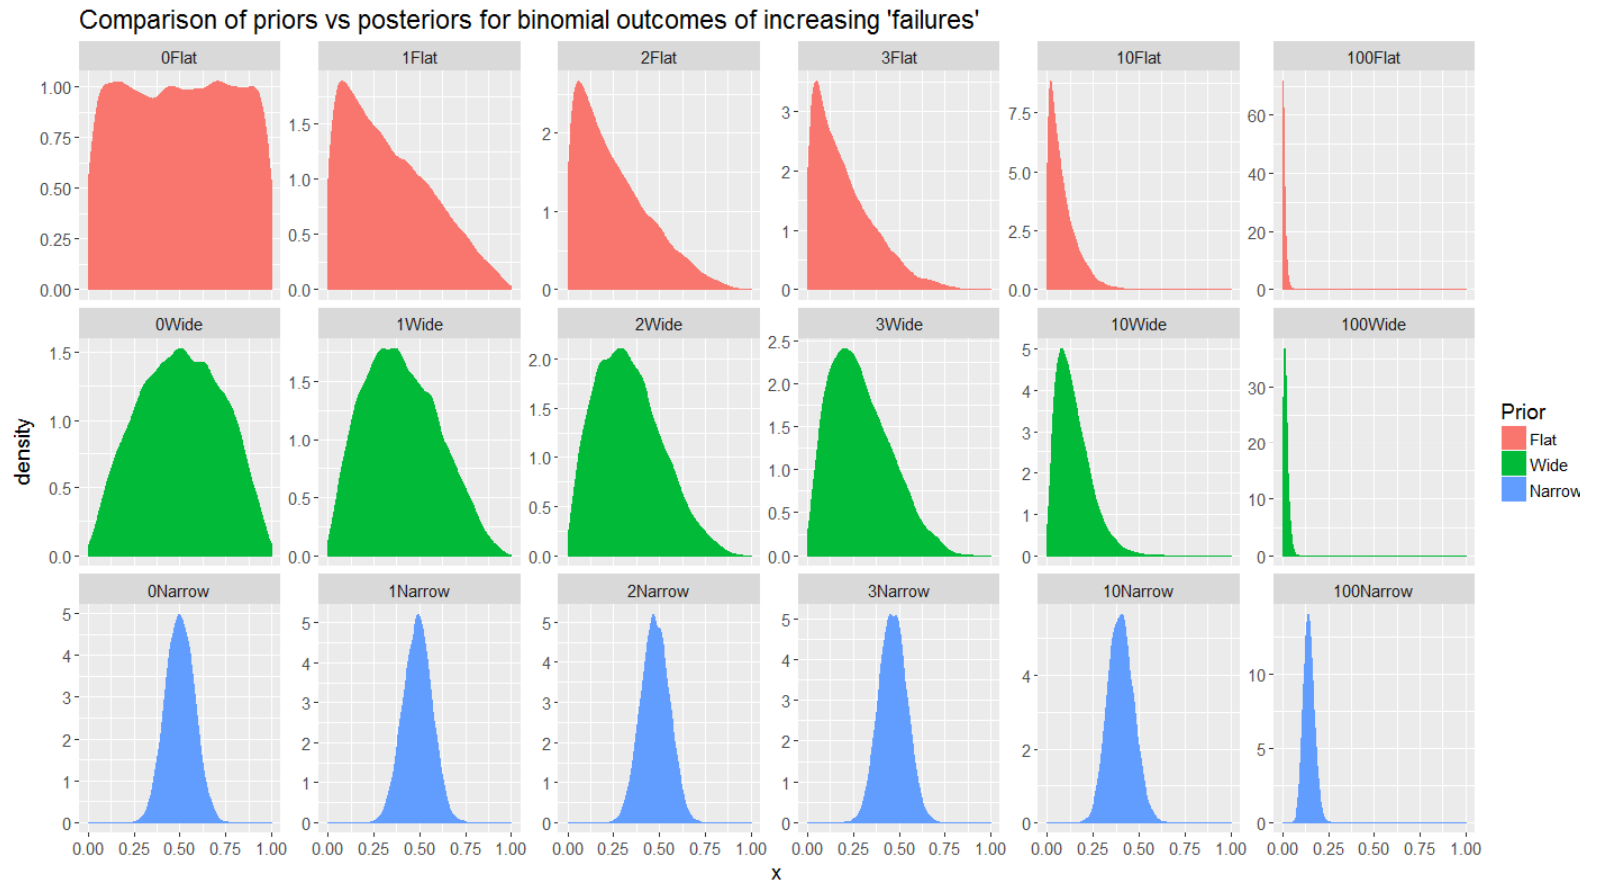
$$p(y) = \sum_{\theta} p(\theta)p(y|\theta)$$

$$\int p(\theta)p(y|\theta)d\theta \quad (\text{continuous})$$

> priors

noninformative vs
informative

conjugacy



Model:

$Y \sim \text{normal}(\mu, \sigma);$

$\mu = a + b \cdot X;$

Priors:

$\sigma \sim \text{cauchy}(0, 5);$


$a \sim \text{normal}(0, 10);$

$b \sim \text{normal}(0, 10);$

- Numeric integration (MCMC)
- Multiple chains

> MCMC

- Numeric integration
- Multiple chains



> Why?

- Philosophical reasons
- Be explicit about prior info / state of knowledge
- Decision making (optimisation, utility, etc)
- Updating
- Model structures
- Characterise and propagate uncertainty
- Combine (causal) inference, prediction and assessment

rstan> installation

- Stan is a probabilistic programming language
- C++ library for Bayesian modeling and inference
- Can be run from R (rstan)
- Primarily uses the No-U-Turn sampler (NUTS)
- Needs rtools
- `install.packages("rstan", dependencies=TRUE)`

```
devtools::find_rtools()
```

```
Sys.getenv('PATH')
```

<https://github.com/stan-dev/rstan/wiki/Installing-RStan-on-Windows>

> workflow

(write model)

- Data (list)
- Model (save as .stan file best)
- Compile model
- Run model (pass data and model spec. to stan)
- Analyse output and 'extract' samples.

(Compare models etc.)

> model.blocks

```
functions {  
}  
data {  
}  
transformed data {  
}  
parameters {  
}  
transformed parameters {  
}  
model {  
}  
generated quantities {  
}
```

- Line must end in ";"
- Comments "//"
- Models support vectors



> `baisc.model(dnorm)`

```
data{
  int<lower=0>      N; // No. obs
  vector<lower=0>[N] x; // data
}
parameters{
  real              mu; // mean
  real<lower=0>    sigma; // standard deviation
}
model{
  x ~ normal(mu, sigma); // univariate distribution
}
```

Inference for Stan model: 85831eef6fda8921202a26d9d49640ee.
3 chains, each with iter=100; warmup=50; thin=1;
post-warmup draws per chain=50, total post-warmup draws=150.

	mean	se_mean	sd	2.5%	25%	50%	75%	97.5%	n_eff	Rhat
mu	0.50	0.00	0.00	0.49	0.50	0.50	0.50	0.50	123	1.01
sigma	0.10	0.00	0.00	0.10	0.10	0.10	0.10	0.10	38	1.02
lp__	1802.66	0.15	1.06	1800.11	1802.49	1803.04	1803.28	1803.43	51	1.06

Samples were drawn using NUTS(diag_e) at Mon Sep 11 16:22:37 2017.
For each parameter, n_eff is a crude measure of effective sample size,
and Rhat is the potential scale reduction factor on split chains (at
convergence, Rhat=1).

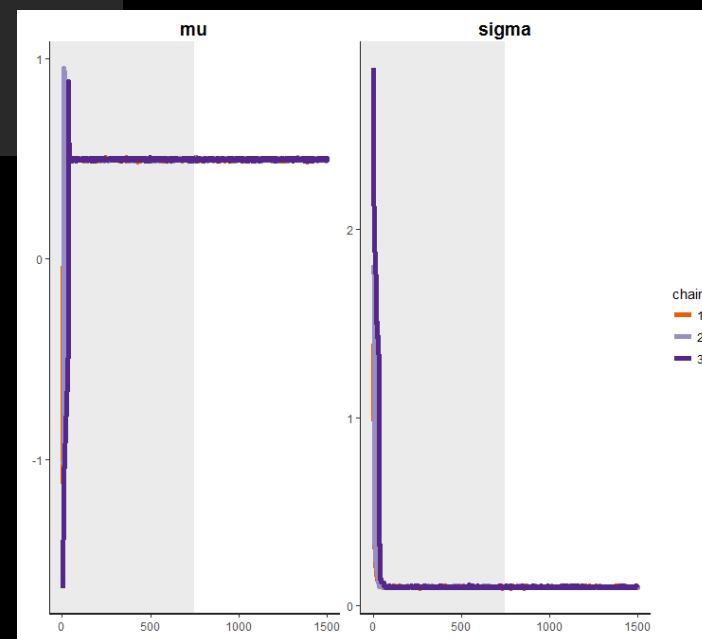



More samples: iter=1500


Inference for Stan model: 85831eef6fda8921202a26d9d49640ee.
3 chains, each with iter=1500; warmup=750; thin=1;
post-warmup draws per chain=750, total post-warmup draws=2250.

	mean	se_mean	sd	2.5%	25%	50%	75%	97.5%	n_eff	Rhat
mu	0.50	0.00	0.00	0.49	0.50	0.50	0.50	0.50	2250	1
sigma	0.10	0.00	0.00	0.10	0.10	0.10	0.10	0.10	1204	1
lp__	1802.52	0.03	0.89	1800.16	1802.13	1802.77	1803.16	1803.42	1227	1

Samples were drawn using NUTS(diag_e) at Mon Sep 11 16:33:31 2017.
For each parameter, n_eff is a crude measure of effective sample size,
and Rhat is the potential scale reduction factor on split chains (at
convergence, Rhat=1).



```
>model.checking
```




```
> baisc.model(dnorm)
```

```
model{
```

```
    x ~ normal(mu, sigma);           // vector notation
```

```
    target+= normal_lpdf(x | mu2, sigma2); // increment
```

```
for(i in 1:N){
```

```
    x[i] ~ normal(mu3, sigma3); // loop style
```

```
    }
```

```
}
```

> linear.model

```
data {  
  int<lower=0> N; // Num obs  
  vector[N] Y;   // Obs times  
  vector[N] X;   // Obs Xances  
}
```

```
parameters {  
  real alpha;  
  real beta;  
  real<lower=0> sigma;  
}  
model {  
  // priors -----  
  alpha ~ normal(0,10);  
  beta ~ normal(0,10);  
  sigma ~ cauchy(0,5);  
  // model -----  
  Y ~ normal(alpha + beta * X, sigma);  
}
```

> linear.model

3 chains, each with iter=1000; warmup=500; thin=1;
post-warmup draws per chain=500, total post-warmup draws=1500.

	mean	se_mean	sd	2.5%	25%	50%	75%	97.5%	n_eff	Rhat
alpha	-2.47	0.04	1.00	-4.49	-3.04	-2.46	-1.85	-0.56	514	1.01
beta	2.96	0.00	0.06	2.86	2.93	2.96	2.99	3.07	534	1.01
sigma	1.21	0.02	0.39	0.74	0.96	1.14	1.35	2.25	494	1.01

```
> runf <- run[1:10,] %>% select(Distance, Minutes)
> fit <- lm(Minutes~Distance, data=runf)
> fit
```

Call:


```
lm(formula = Minutes ~ Distance, data = runf)
```

Coefficients:

(Intercept)	Distance
-2.488	2.963

> linear.model

Running records



> log.rules

- Calculations often involve product of small probabilities
- Smaller than machine epsilon/precision
- Use logs to avoid under/overflow and floating point errors:

$$\log(a*b) == \log(a) + \log(b)$$

$$\log(a/b) == \log(a) - \log(b)$$

>loo(“leave one out”)

- approximate LOO using Pareto Smoothed Importance Sampling (PSIS).

elpd_loo - (expected log predictive density),

p_loo - (effective number of parameters), and

looic - (the LOO information criterion).

- `compare(model1, model2)`
- WAIC, DIC, BIC etc.
- Model comparison/averaging

<https://github.com/stan-dev/loo/blob/master/vignettes/loo-example.Rmd>

> predictive.evaluation(lppd)

*Out of Sample predictive evaluation

generated quantities {

```
real log_p_new;      // posterior predictive log density test data
```

```
  log_p_new = 0;
```

```
  for(i in 1:N_nu){
```

```
log_p_new = log_p_new + normal_lpdf(Y_nu[i] | alpha + X_nu[i]*beta, sigma);
```

```
}}
```

- NB: this gives posterior expectation of the log predictive density:

(expectation of the log)

- We want: log of posterior predictive density (lppd):

(log of the expectation)

Therefore needs to be 'corrected'

$(-\log(m) + \log_{\text{sum_exp}}(\log p(y_{\text{nu}} | \theta)))$

<http://mc-stan.org/users/documentation/case-studies/pool-binary-trials.html> (also McElreath, p193)

> stan.advantages

- Support, resources / momentum
- Development
- algorithms (HMC/NUTS)
- faster convergence & better explore parameter space
- vectorisation
- supports parallelisation
- functions, distributions
- saving compiled models

> bayes.challenges

- Complicated
- Time consuming
- Unresolved issues / developing field
- thinning, h-centering, priors, goodness-of-fit etc.
- Computational complexity / large models
- (Latent) discrete unknown parameters (stan – can be done but not directly)

> other.packages

- shinystan
- ggmmcmc
- loo
- rstanarm*
- brms*
- Rethinking (github)*
- Prophet

*Do some of the model building for you

> resources.bayes

Andrew Gelman, John B. Carlin, Hal S. Stern, David B. Dunson, Aki Vehtari, Donald B. Rubin, 2013, **Bayesian Data Analysis, Third Edition**, CRC press.

John Kruschke, 2014, **Doing Bayesian Data Analysis, 2nd Edition**: A Tutorial with R, JAGS, and Stan, Elsevier.

Richard McElreath, 2015, **Statistical Rethinking**: A Bayesian Course with Examples in R and Stan, CRC press.

Adam Branscum, Ronald Christensen, Timothy E Hanson, and Wesley O. Johnson, 2010, **Bayesian Ideas and Data Analysis**: An Introduction for Scientists and Statisticians, CRC press

John Kruschke & Liddell, T., 2017, **Bayesian data analysis for newcomers**, *Psychonomic Bulletin & Review*, <https://psyarxiv.com/nqfr5>

Philosophy:

John K. Kruschke and Liddell, T.M., 2017, **The Bayesian New Statistics**: Hypothesis testing, estimation, meta-analysis, and power analysis from a Bayesian perspective, *Psychonomic Bulletin & Review*

Andrew Gelman and Shalizi, R.S., 2012, **Philosophy and the practice of Bayesian statistics**, *British Journal of Mathematical and Statistical Psychology*, 66:8-38

> resources.stan

Stan reference manual

<http://mc-stan.org/>

<http://mc-stan.org/users/documentation/index.html>

<http://discourse.mc-stan.org/>

<https://github.com/stan-dev/rstan>

<https://cran.r-project.org/web/packages/rstan/vignettes/rstan.html>

<http://modernstatisticalworkflow.blogspot.com.au/>

<http://doingbayesiandataanalysis.blogspot.com>

<http://www.mcmchandbook.net/>

Kaggle:

<http://blog.kaggle.com/2017/05/19/march-machine-learning-mania-1st-place-winners-interview-andrew-landgraf/>